

BASIC DATABASES

View, Stored Procedure, Index

NGUYEN Hoang Ha


Email: nguyen-hoang.ha@usth.edu.vn



VIEW

View concept

- A view is a “virtual” or logical table that is derived from other tables



Pros vs. Cons

■ Pros:

- Simplify complex queries
- Enable computed columns
- Provide a security layer: hide sensitive data
- Enable backward capability

■ Cons

- Performance
- Table dependency: table changes → need to change views

Syntax

```
CREATE [ALGORITHM = {MERGE | TEMPTABLE | UNDEFINED}]  
VIEW view_name [(column_list)]  
AS  
select-statement [WITH CHECK OPTION];
```

■ ALGORITHM

- **MERGE:** MySQL combines input query with the select-statement. MERGE is not allowed if the SELECT statement contains aggregate functions or `DISTINCT`, `GROUP BY`, `HAVING`, `LIMIT`, `UNION`, `UNION ALL`, subquery, SELECT statement refers to no table.
- **TEMPTABLE:** MySQL creates a temporary table based on the SELECT statement that defines the view, then performs query against this temporary table.
- **UNDEFINED:** MySQL makes choice of MERGE or TEMPTABLE.

View examples

■ Computed columns

```
CREATE VIEW sale_per_order AS
SELECT order_id, SUM(quantity * unit_price * (1-discount)) total
FROM order_details
GROUP BY order_id
ORDER BY total DESC;
```

■ Based on a sub query

```
CREATE VIEW above_avg_products AS
SELECT product_code, product_name, list_price
FROM products
WHERE list_price > (SELECT AVG(list_price)
FROM products)
ORDER BY list_price DESC;
```

■ Based on another view

```
CREATE VIEW big_sale_orders AS
SELECT order_id, ROUND(total,2) AS total
FROM sale_per_order
WHERE total > 1000;
```

Updatable views

- SELECT statement defining the view must not contain following elements:
 - Aggregate functions such as MIN, MAX, SUM, AVG, and COUNT.
 - DISTINCT
 - GROUP BY clause.
 - HAVING clause.
 - UNION or UNION ALL clause.
 - Left join or outer join.
 - Subquery in the SELECT clause or in the WHERE clause that refers to the table appeared in the FROM clause.
 - Reference to non-updatable view in the FROM clause.
 - Reference only to literal values.
 - Multiple references to any column of the base table

WITH CHECK OPTION Clause

- Role: to prevent updating or inserting rows that are not visible through the view
- Example

```
CREATE OR REPLACE VIEW northwind_products
AS
SELECT id, product_code, product_name
FROM products
WHERE product_name LIKE 'Northwind%'
WITH CHECK OPTION;
```

```
INSERT INTO northwind_products (product_code, product_name)
VALUES ('HNB', 'Hanoi Beer'); -- This is invalid
```

```
INSERT INTO northwind_products (product_code, product_name)
VALUES ('NWnew', 'Northwind Beer');
```

```
UPDATE northwind_products
SET product_name = 'Nwd beer'
WHERE product_code = 'NWnew'; --- WITH CHECK OPTION will prevent this statement from running
```


View management

- Show view definition

- `SHOW CREATE VIEW [database_name].[view_name];`

- Delete view:

- `DROP VIEW [IF EXISTS] view_name`

- Change view

- `ALTER[ALGORITHM = {MERGE | TEMPTABLE | UNDEFINED}]
VIEW view_name [(column_list)]
AS
select-statement [WITH CHECK OPTION];`

- OR: CREATE OR REPLACE VIEW

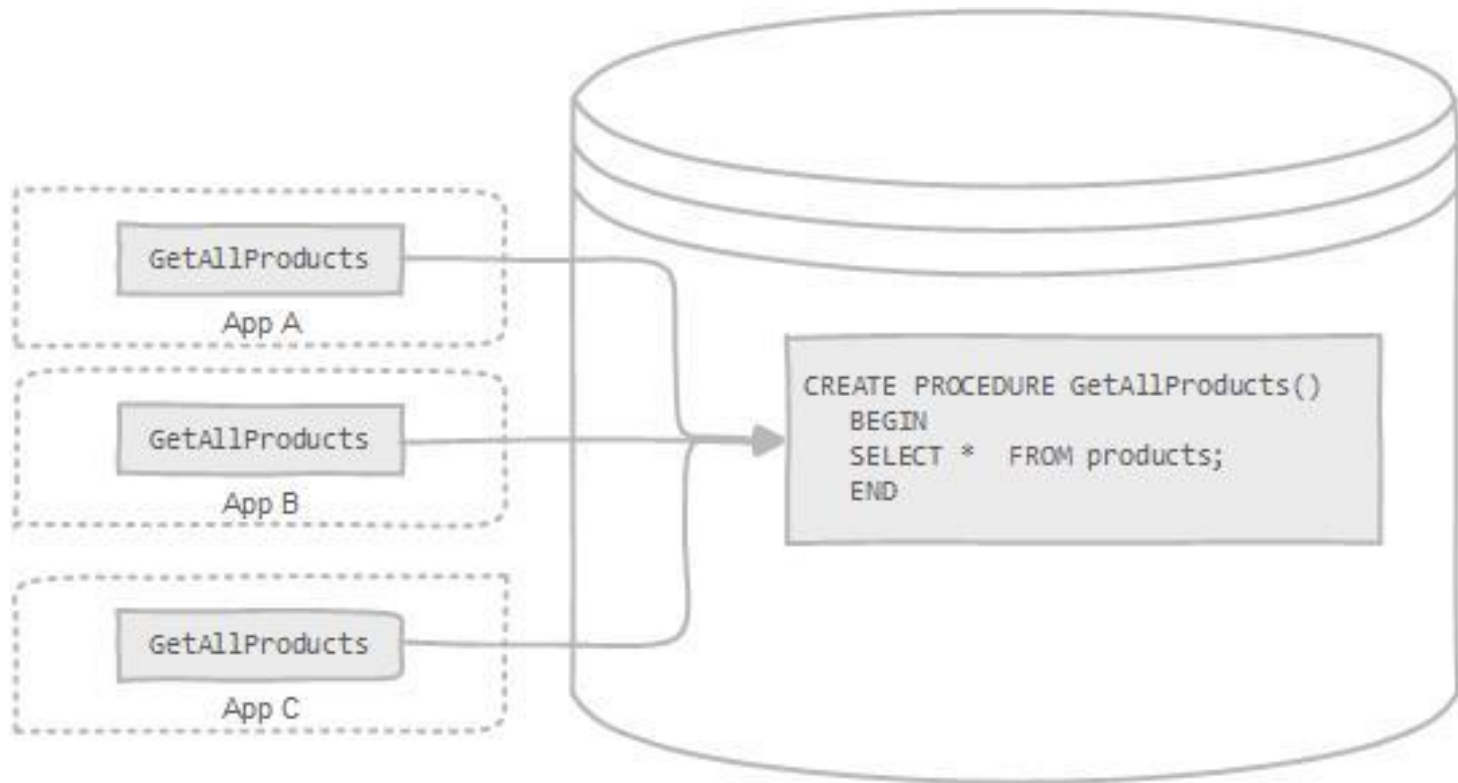
- List all views with updateable information

- `SELECT table_name, is_updatable
FROM information_schema.views`

STORED PROCEDURE

Concept

- A stored procedure is a segment of SQL statements stored inside the database catalog



Pros vs. Cons

- Pros:

- Better performance
- Reduce traffic
- Be reusable and transparent
- Provide a secure way to access data

- Cons

- CPU usage can increase if logical operators are overused
- Hard to debug, maintain

Example

```
DROP PROCEDURE IF EXISTS count_products;
delimiter //
CREATE PROCEDURE count_products (OUT param1 INT)
BEGIN
    SELECT COUNT(*) INTO param1
    FROM products;
END//
delimiter ;

CALL count_products (@a); SELECT @a;
```

Input parameter

```
DELIMITER //  
CREATE PROCEDURE get_customers_by_city(IN search_city nvarchar(255))  
AS  
BEGIN  
    SELECT * FROM customers  
    WHERE city = search_city;  
END //  
DELIMITER ;
```

```
CALL get_customers_by_city('Seattle');
```

INDEX

Why indexing?

- Indexing are one of the most important and useful tools for achieving high performance in a relational database
- Many database administrators consider indexes to be the single most critical tool for improving database performance
- An index is a data structure that contains a copy of some of the data from one or more existing database tables
- A database index provides an organizational framework that the DBMS can use to quickly locate the information that it needs
- This can vastly improve the speed with which SQL queries can be answered

Without index

- Query for a random name within the table
- What is the average search time if the process is repeated many times?

$$average = \frac{n + 1}{2}$$

- What is the maximum search time?

$$Maximum = n$$

Row Position	Last Name
6	Al Rabeeah
16	Beena
13	Doshi
10	Flores
11	Fung
19	Gani
8	Garcia
21	Hu
22	Israr
9	Johnson
18	Ly
2	Mishra
17	Ngo
20	Pham
3	Salehian
1	Schluter
14	Scruton
12	Spievak
4	Vu
5	Wah
15	Winter
7	Wong

With index

- Query for a random name within the table
- What is the average search time if the process is repeated many times?

$$average = \log_2(n) - 1 = 3.5$$

- What is the maximum search time?

$$Maximum = \log_2(n) = 4.5$$

Row Position	Last Name
1	Schluter
2	Mishra
3	Salehian
4	Vu
5	Wah
6	Al Rabeeah
7	Wong
8	Garcia
9	Johnson
10	Flores
11	Fung
12	Spievak
13	Doshi
14	Scruton
15	Winter
16	Beena
17	Ngo
18	Ly
19	Gani
20	Pham
21	Hu
22	Israr

Index concepts

- Indexes are created on one or more columns in a table
 - For example:
 - An index is created on a PK column
 - The index will contain the PK value for each row in the table, along with each row's ordinal position (row number) within the table
 - When a query involving the PK is run, the DBMS will find the PK value within the index. The DBMS will then know the position of the row within the table
 - The DBMS can then quickly locate the row in the table that is associated with the PK value
- Without an index, the DBMS has to perform a table scan in order to locate the desired row

Index concepts

- An index can be created on most, but not all, columns. Whether an index can be created on a column depends on the column's datatype
- Columns with large object data types cannot be indexed without employing additional mechanisms These data types include:
 - Text
 - ntext
 - Image
 - varchar (max)
 - Nvarchar(max)
 - varbinary(max)

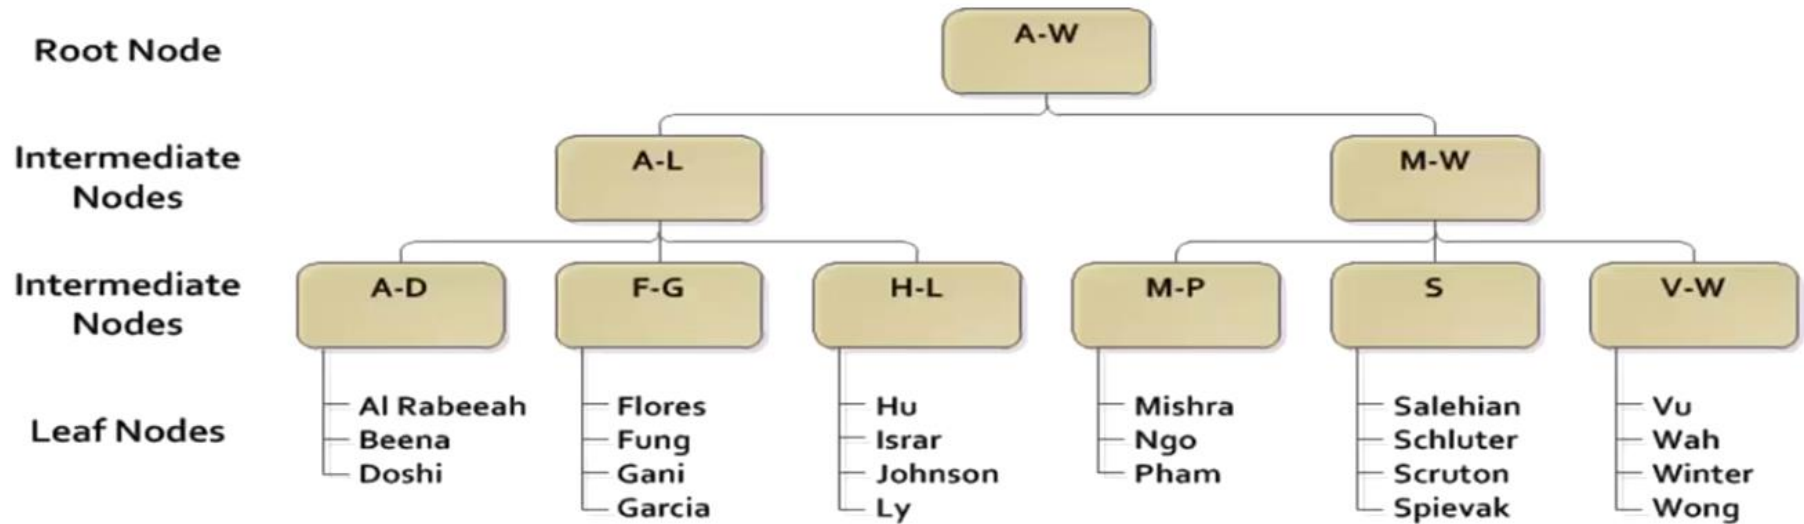
Index concepts

- Creating an index increases the amount of storage space required by the database
 - This occurs because an index contains a copy of some of the data in a table
 - To estimate the storage space requirements of an index, we can use the following formula:
- $\text{Number of rows in table} \times \text{Average number of bytes required per row for the indexed columns}$

B-Tree Index

- Balance-Tree: the most common type of database indexing
- B-trees use pointers and several layers of nodes in order to quickly locate desired data
- Root node
- Intermediate nodes
- Leaf nodes
- When the DBMS processes a query which includes an indexed column, it starts at the root node of the B-tree and navigates downward until it finds the desired leaf

B-tree example



Clustered Indexes

- In a clustered index, the actual data rows that comprise the table are stored at the leaf level of the index
- The indexed values are stored in a sorted order
 - This means that there can be only one clustered index per table
 - PK columns are good candidates for clustered indexes

Clustered B-Tree Example

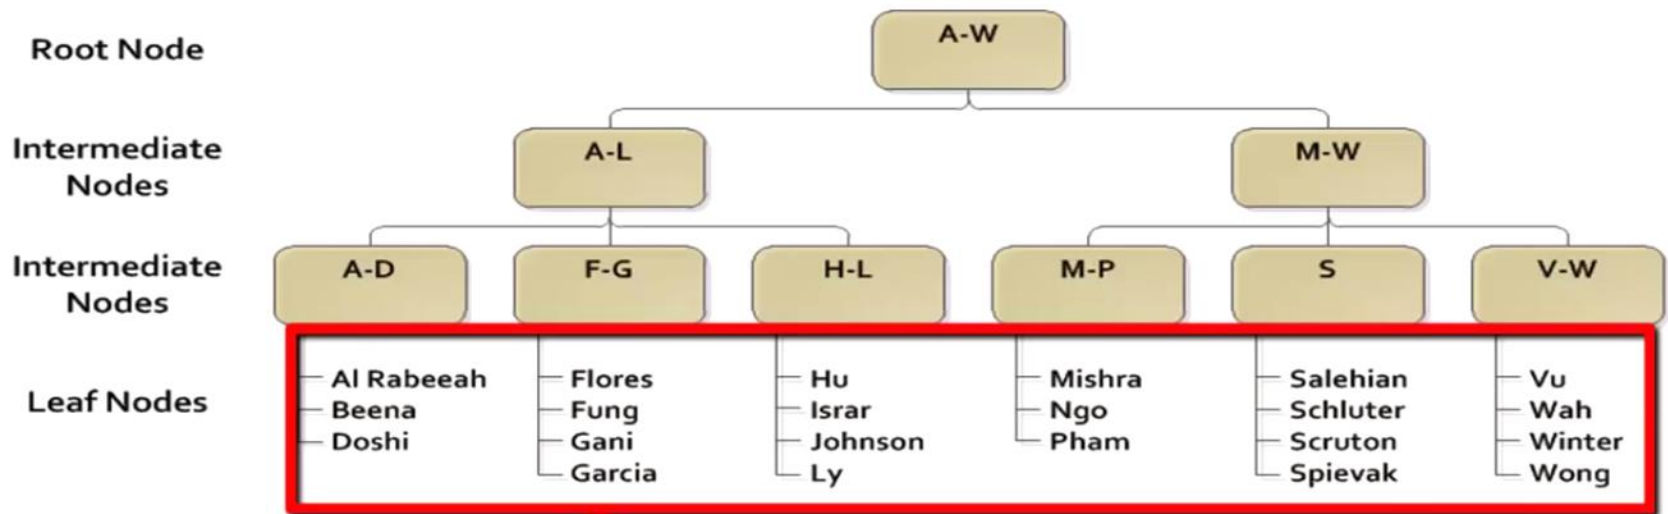


table rows instead of pointers

Indexing Guidelines

- If a table is heavily updated, index as few columns as possible
- If a table is updated rarely, use as many indexed columns as necessary to achieve maximum query performance
- Clustered indexes are best used on columns that do not allow null values and whose value are unique
- The performance benefits on an index are related to the uniqueness of the values in the indexed column
 - Index performance is poor when an indexed column contains a large proportion of duplicate values
 - Index performance is best when an indexed column contains unique values